




Politecnico
di Torino



Exercises

Advanced Machine Learning

Teaching Assistant:


Gabriele Tiboni

gabriele.tiboni@polito.it

Overview

1. Train a **Convolution Neural Network** for image classification:
 - CNN: **AlexNet**
 - Images: **Caltech-101**
2. Exercise **Steps**:
 - Before you start
 - Data preparation
 - Train from scratch
 - Transfer Learning
 - Data Augmentation
 - Beyond AlexNet

AlexNet ([link](#))



Training resources

Deep Learning requires **GPU** acceleration 

Training resources

Deep Learning requires **GPU** acceleration 

You can use [Google Colab](#)! 

colab



PYTORCH

Code templates

1. The template of the main code is available [here](#):
 - [Save a copy on your own drive](#): “File” -> “Save a copy in Drive”
 - Switch to GPU acceleration: “Edit” -> “Notebook Settings”
2. Caltech-101 and the dataset class template are available via [this](#) GitHub repository.

Step 0: Before you start

Before you start

1. Study code and data:
 - a. Read carefully the template code (including the comments) to understand how everything is done. Pay extra attention to: Preprocessing, Datasets, Training, Testing
 - b. Explore the data provided on the GitHub repository.

2. Run the code:
 - a. training should take less than 10 minutes
 - b. you have to stay connected

Step 1: Data Preparation

Step 1: Data Preparation

1. Create your own dataset class for Caltech-101, following the code provided in the GitHub repository (`caltech_dataset.py`):
 - a. train-test split is already provided in the github repository: `test.txt` and `train.txt` contain the relative paths for all the images they include
 - b. there is also a `BACKGROUND` folder that you are required to filter out
 - c. the class should read and store only the images belonging to the corresponding split (see the `split` parameter of the Caltech class)

Step 2: Train from scratch

Step 2: Train from scratch

During training, we need a **validation set** for hyperparameter tuning and model selection.

In deep learning, it is often **prohibitive** to perform a **K-fold Cross Validation** since training takes a lot of time.

The most common procedure is the hold-out, having one single validation set both for hyperparameter tuning and model selection.

Step 2: Train from scratch

1. **Split** the training set in **training** and **validation sets**:
 - a. The validation set must be half the size of the training set
 - b. Be careful to not filter out entire classes from the sets!
 - c. Train and validation must be balanced: aim for half samples of each class in train and the other half in val
2. Implement **model selection** with validation:
 - a. After each training epoch, evaluate (= test) the model on the validation set
 - b. Use only the best performing model on the validation set for test

Step 2: Train from scratch

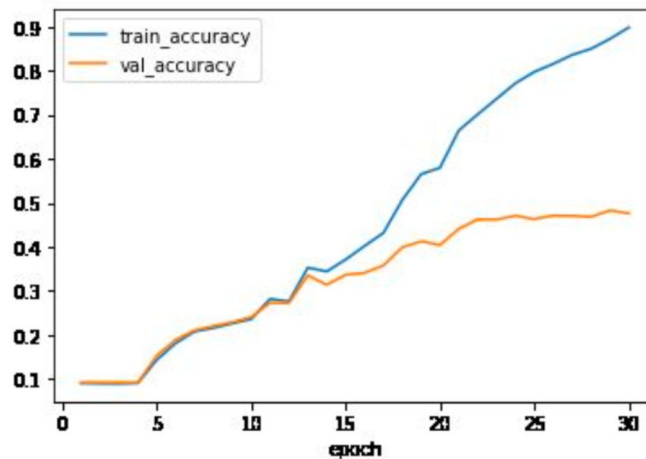

The current implementation trains using SGD with momentum for 30 epochs with an initial learning rate (**LR**) of 0.001 and a decaying policy (STEP_SIZE) after 20 epochs.

3. Experiment with **at least two different sets of hyperparameters**:

- a. The first hyperparameter to optimize is the learning rate
- b. Low learning rate = the model converges too slowly (loss decreases slowly)
- c. High learning rate = the model converges fast but accuracy is sub-optimal
- d. Too high learning rate = the model diverges (loss increases)

suggestion: experiment changing LR and one among: decaying policy, optimiser, epochs)

Step 2: Train from scratch



Step 3: Transfer Learning

Step 3: Transfer Learning

Deep Learning needs very large datasets to train good features but **Caltech-101 is very small dataset**.

Solution 1: Transfer Learning

- Use the weights learned by training on a large related dataset as a starting point for training on the small dataset
- We can also choose to “freeze” part of the network and only train certain layers

Step 3: Transfer Learning

Solution 1: Transfer Learning

1. Load AlexNet with weights trained on the **ImageNet** dataset.
2. Change the **Normalize** function of Data Preprocessing to Normalize using ImageNet's mean and standard deviation
3. Run experiments with at least **three different sets** of hyperparameters
4. Experiment by training only the fully connected layers (freeze other layers)
5. Experiment by training only the convolutional layers:
 - a. Compare all results you obtained, reasoning about what does it mean logically to freeze some layers

Step 4: Data Augmentation

Step 4: Data Augmentation

Deep Learning needs very large datasets to train good features but **Caltech-101 is very small dataset.**

Solution 2: Data Augmentation

- Artificially increase the dataset size by applying some transformations at training time, preserving the label
- Transformations applied should be the ones we are expected to see at test time, otherwise the overall accuracy could be negatively affected

Step 4: Data Augmentation

Solution 2: Data Augmentation

1. Apply at least three different sets of preprocessing for training images:
 - a. [link](#)
 - b. no need to use complex transformations, try with the simple ones, such as **flipping** or changing the **brightness** and combine them
 - c. if test data is very similar to training data, some data augmentation policies may worsen accuracy
 - d. if the loss keeps decreasing, increase training epochs or increase learning rate
2. Compare the results

Step 5: Beyond AlexNet

Step 5: Beyond AlexNet

1. Try with different models (e.g. **VGG**, **ResNet**)
 - a. check their requirements in terms of:
 - i. input image size
 - ii. GPU memory consumed (you may need to significantly reduce batch size)
2. Compare the results

Now it's your turn, try!

